

# What is Machine Learning?

Oxford Spring School in Advanced Research Methods, 2021

Dr Thomas Robinson, Durham University

26/03/2021

# Introduction

# This course

5  $\times$  3 hour sessions, covering:

1. Introduction to ML and maximum likelihood estimation
2. ML extensions to regression
3. Tree-based methods
4. Neural networks
5. Ensemble methods

The logic:

- ▶ Understand the underlying mechanics of parameter estimation
- ▶ Starting from a modeling strategy we are likely familiar with. . .
- ▶ . . . moving on to those that are algorithmically more complicated
- ▶ Building on the same foundational concepts across each day

# Balancing a course on machine learning

- ▶ It's possible to build a five-day course on any one of the five topics we will cover
- ▶ And there are lots of topics we will not cover:
  - ▶ Unsupervised methods
  - ▶ Clustering algorithms
  - ▶ Text-specific ML models

Goal of this course is to provide the *fundamentals* that can be applied across ML contexts

- ▶ Emphasis on accessibility
- ▶ Transferability
- ▶ Relevance to our social science research streams

# Session structure

Each session will be a mixture of:

- ▶ Lecture content and discussion
  - ▶ I will leave plenty of opportunities for questions
  - ▶ Building in some time for us to think through some applied problems
- ▶ Coding walkthroughs (approx. 1 hour)
  - ▶ Hands on experience using some of the algorithms we'll study
  - ▶ Many of these are now very easy to implement

# Today's session

Goals are threefold:

1. Introduce the topic of machine learning
  - ▶ What is ML?
  - ▶ How do we distinguish it from statistics?
  - ▶ What sort of problems might we apply it to?
2. Introduce key conceptual distinctions we will make throughout the course
3. Introduce maximum likelihood estimation
  - ▶ A key way in which ML parameters are estimated
  - ▶ Build our own logistic regression estimator

What is machine learning?

# (Machine) learning and statistics

ML is a vague field:

*“Machine learning is a subfield of computer science that is concerned with building algorithms which, to be useful, rely on a collection of examples of some phenomenon. . . the process of solving a practical problem by 1) gathering a dataset, and 2) algorithmically building a statistical model based on that dataset.” – Burkov 2019*

*“Statistical learning refers to a set of tools for modeling and understanding complex datasets. It is a recently developed area in statistics and blends with parallel developments in computer science and, in particular, machine learning” – James et al. 2013*



# Prediction and inference

ML typically focuses on prediction problems instead of inference problems:

- ▶ Inference is concerned with estimating the size/direction of the relationship between variables ( $\hat{\beta}$  problems)
- ▶ Prediction is concerned with estimating some outcome, using the relationships between variables ( $\hat{y}$  problems)

These two facets are clearly connected:

- ▶ If we know the size/direction of the relationships, we can predict the outcome
- ▶  $\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_{1i} + e_i$
- ▶ But we rarely know (or even pretend to know) the true model
- ▶ This uncertainty means we need statistics/ML for both inference *and* prediction

# Classification and prediction

We can distinguish two types of prediction problem [CITE VARIAN]

- ▶ **Prediction** – estimating the value of a continuous variable (sometimes referred to as “regression” problems)
  - ▶ E.g. The ideology of a politician in 2D space
  - ▶ The number of votes received by a candidate
  - ▶ The “trustworthiness” of an individual [CITE NATURE PAPER]
- ▶ **Classification** – estimating which *class* of a category an observation belongs to
  - ▶ Party identity (Republican/Democrat/Libertarian/Independent)
  - ▶ The topic of a tweet (foreign/domestic, pro/con)

## $\hat{Y}$ and $\hat{X}$ problems

We can also think about where the prediction problem lies:

- ▶  $\hat{y}$  problems are about the dependent variable
  - ▶ To predict an election winner...
  - ▶ ... or the probability of revolution...
  - ▶ ... or the weather tomorrow
  - ▶ These are not necessarily inferential problems
- ▶  $\hat{X}$  problems are about independent variables
  - ▶ Dimensions of interest that may be important to our theory...
  - ▶ ... but which are not directly observable (i.e. latent)
  - ▶ We want to make predictions over  $\mathbf{X}$  so we can test an inferential theory about the relationship between  $X$  and  $y$

# Supervised vs unsupervised methods

## Supervised methods

- ▶ Contains labeled examples of observations with corresponding outcomes  $\mathbf{y}$  – the **training data**
- ▶ Use these examples to “learn” the relationship between  $\mathbf{y}$  and  $\mathbf{X}$
- ▶ Then predict  $\mathbf{y}$  for a *new* **test** dataset  $\mathbf{X}^{\text{TEST}}$  where  $\mathbf{y}^{\text{TEST}}$  is not observed

Learning the relationship:

$$\underbrace{\begin{bmatrix} 1 \\ 0 \\ 0 \\ \vdots \\ 1 \end{bmatrix}}_{\mathbf{y}^{\text{TRAIN}}} \quad \underbrace{\begin{bmatrix} 3.3 & 1.1 & 0 \\ 2.7 & 0.8 & 0 \\ 1.8 & 0.1 & 1 \\ \vdots & \vdots & \vdots \\ 5 & 1.2 & 0 \end{bmatrix}}_{\mathbf{X}^{\text{TRAIN}}}$$

Predicting on new data

$$\mathbf{X}^{\text{TEST}} = \begin{bmatrix} 3.5 & 1.9 & 1 \\ 5.4 & 0.3 & 0 \\ 1.7 & 0.5 & 1 \end{bmatrix}$$

# Machine learning

Expectation: I need a \$1m super computer

Reality: It runs in minutes on a personal computer



Figure 1: Google: Tensor Processing Unit server rack

# Why machine learning?

Machine learning can be:

- ▶ Powerful
- ▶ Flexible
- ▶ Reduce the burden on the researcher

It helps solve lots of **prediction problems**...

... and can assist in **inference problems** too

- ▶ Through  $\hat{X}$  and  $\hat{Y}$  problems

# Machine learning and social science

But ML is not a panacea!

- ▶ ML cannot solve problems of poor research design
- ▶ And can introduce it's own issues

## Twitter apologises for 'racist' image-cropping algorithm

**Users highlight examples of feature automatically focusing on white faces over black ones**

## Minority Report-style tech used to predict crime is 'prejudiced'

In his first interview since becoming surveillance commissioner, Fraser Sampson warns about accuracy of predictive policing technology

## Maximum Likelihood Estimation (a gentle introduction)



# Notation

Throughout the course, we will follow the notation set out in Buryiv (2019):

- ▶  $\theta$  is a scalar - i.e. a single number
  - ▶ E.g.,  $\theta = 3.141$ ,  $x = 1$  etc.
- ▶  $\boldsymbol{\theta}$  is a vector - i.e. an ordered list of scalar values
  - ▶ E.g.,  $\boldsymbol{\theta} = [0.5, 3, 2]$
- ▶  $\mathbf{X}$  is a matrix
  - ▶ E.g.,

$$\mathbf{X} = \begin{bmatrix} 1 & 5 \\ 24 & -3 \end{bmatrix}$$

- ▶  $\mathbf{x}^{(k)}$  is the  $k$ th column of  $\mathbf{X}$
- ▶  $\mathbf{x}_i$  is the  $i$ th row in matrix  $\mathbf{X}$
- ▶  $x_i^{(k)}$  is the  $k$ th element of the row vector  $\mathbf{x}_i$

# Probability notation

Let  $p$  denote a probability distribution *function* that returns the probability of an event or observation:

- ▶  $p(A) = 0.5$  means the probability of event  $A$  is 0.5
- ▶  $0 \leq p(A) \leq 1$
- ▶ Conditional probability  $p(A|c) = 0.25$  means the probability of  $A$  **given** (or conditional on)  $c$  is 0.25.
- ▶  $p(A \text{ and } B) = p(A)p(B|A)$
- ▶ If  $p(A) \perp p(B)$ , then  $p(B|A) = p(B)$ 
  - ▶ So, if  $p(A) \perp p(B)$ , then  $p(A \text{ and } B) = p(A)p(B)$

# Notation quiz

What are the following:

1.  $\mathbf{a}$
2.  $y_i$
3.  $\beta$
4.  $\beta$
5.  $\ominus$

If  $p(A) = 0.5$ ,  $p(B) = 0.1$ , and  $p(B|A) = 0.3$ :

6. Is  $p(A) \perp p(B)$ ?
7. What is  $p(A \text{ and } B)$ ?

## Bayes Theorem (from a frequentist perspective)

$$\underbrace{P(A|B)}_{\text{Posterior}} = \frac{\overbrace{P(B|A)}^{\text{Likelihood}} \times \overbrace{P(A)}^{\text{Prior}}}{\underbrace{P(B)}_{\text{Evidence}}}$$

We can use Bayes formula to estimate the posterior probability of some parameter  $\theta$ :

$$p(\theta|\mathbf{X}) \propto p(\mathbf{X}|\theta) \times p(\theta),$$

where  $\mathbf{X}$  is the data.

## Likelihood function

Let's suppose that we have no prior knowledge over  $\theta$ , so we'll drop the prior and focus specifically on the likelihood:

$$\mathcal{L}(\theta) = p(\mathbf{X}|\theta)$$

*How would we calculate this?*

$$\begin{aligned}\mathcal{L}(\theta) &= p(\mathbf{x}_1|\theta) \times p(\mathbf{x}_2|\theta) \times \dots \times p(\mathbf{x}_n|\theta) \\ &= \prod_{i=1 \dots N} p(\mathbf{x}_i|\theta)\end{aligned}$$

i.e. the product of the probability of each observations within  $\mathbf{X}$ , given  $\theta$ .

*What does this assume?*

## Why is the likelihood function useful?

Suppose we have two alternative values of  $\theta$ :  $\theta_1, \theta_2$ . We can calculate the likelihood *ratio* (LR) of these two possible parameter values:

$$LR = \frac{\mathcal{L}(\theta_1)}{\mathcal{L}(\theta_2)}$$

*If  $LR > 1$ , which parameter value would we pick?*

# Maximum likelihood estimation

We can generalise this for all possible values of  $\theta$ :

$$\arg \max_{\theta \in \Theta} \mathcal{L}(\theta) = \prod_{i=1 \dots N} p(\mathbf{x}_i | \theta)$$

i.e., from the set of all possible parameter values  $\Theta$ , find the parameter value that maximises the likelihood function.

Hence, **maximum** likelihood estimation.

- ▶ How we calculate  $p(\mathbf{x}_i | \theta)$  will depend on the functional form of the underlying distribution
- ▶ We'll explore this specifically with respect to logistic regression later on today

*Why is this useful?*

## Numeric overflow

Multiplying many small numbers means we soon lose the power to calculate them precisely

- ▶ R double-precision numbers range from  $2 \times 10^{-308}$  to  $2 \times 10^{308}$
- ▶ If 400 observations have  $p_{\theta} = 0.01$ ,  $\mathcal{L}(\theta)$  will be outside the computable range

What if we take the log?

- ▶ The log function is strictly increasing
- ▶  $\text{Log}(a \times b) = \text{Log}(a) + \text{Log}(b)$  so we can simply add the values

With the logged likelihood function we do not have the problem of numeric overflow!



# Negative log-likelihood

One final step that we can take is to calculate the *negative* log-likelihood:

- ▶ We then *minimise* the negative log-likelihood
- ▶ We typically want to minimise rather than maximise because many of our procedures for optimisation are based on the former
- ▶ But, broadly, this is just semantics:
  - ▶ Minimising the negative log-likelihood is the same as maximising the log-likelihood

# Logistic regression

Logistic regression:

- ▶ Allows us to estimate  $\beta$  parameters when we have a binary outcome variable
- ▶ More broadly, it is a **binary classification** algorithm – what is the probability that  $y_i = 1$  given a vector of features  $\mathbf{x}_i$ ?

We can write the logistic regression function as,

$$f_{\theta,b}(\mathbf{X}) = \frac{1}{1 + e^{-(\theta\mathbf{X}+b)}}.$$

The goal is to find the *best* values of  $\theta$  and  $b$  that “explains” the data

- ▶ Let's subsume within  $\theta$  s.t.  $\theta = \{b, \theta_1, \dots, \theta_k\}$

# MLE of logistic regression

For a given vector of scalar values  $\theta$ , we can ask what the likelihood of the data is given those values

How do we construct this?

- ▶ If  $y_i = 1$ , we want the  $f_{\theta,b}(\mathbf{x}_i)$
- ▶ But if  $y_i = 0$  we want the inverse, i.e.  $(1 - f_{\theta,b}(\mathbf{x}_i))$
- ▶ We can combine these two using a mathematical “logic gate”:

$$\mathcal{L}_{\theta,b} = f_{\theta}(\mathbf{X})^y \times (1 - f_{\theta}(\mathbf{X}))^{(1-y)},$$

since when  $y_i = 0$ ,  $x^{y_i} = 1$  and  $x^{(1-y_i)} = x$ , and vice versa.

Simplifying, since  $f_{\theta,b}(\mathbf{x}_i) = \hat{y}_i$ :

$$\mathcal{L}_{\theta} = \hat{\mathbf{y}}^y (1 - \hat{\mathbf{y}})^{1-y}$$

## MLE optimization

We can then apply our “tricks” to make the computation easier:


$$-Log(\mathcal{L}_\theta) = -\sum_{i=1}^N Log(\mathcal{L}_\theta(\mathbf{x}_i)),$$

with the goal of minimising this quantity through choosing  $\theta$  and  $b$ .

How exactly do we minimize this function?

- ▶ Unlike OLS, where there is a closed form solution, it is not possible to analytically minimize the negative log-likelihood of the logistic regression
- ▶ We therefore have to use computation to iterate through values of  $\theta$  to approximate the minima

## Minimising the negative log-likelihood in one dimension



# Gradient descent

To find the minimum of the negative log-likelihood we:

1. Choose a value for the starting parameter  $\theta$
2. Calculate the slope of the function at that point
3. Adjust our value of  $\theta$  in the opposite direction to the slope coefficient's sign
4. Recalculate the slope, and repeat 2-4

We can generalise this to  $\theta$ :

- ▶ Let  $Q(\theta)$  be the negative log likelihood function
- ▶ Calculate the **gradient** vector of the function in  $k$ -dimensions
- ▶ Adjust each parameter  $\theta_k \in \theta$  by the negative of the corresponding element of the gradient

$$\theta_k = \theta_k - \lambda \frac{\partial Q(\theta)}{\partial \theta_k}$$

# Logistic regression gradient


The partial derivative for any predictor  $\mathbf{x}^{(j)}$  for the *logistic* cost function is:

$$\frac{\partial Q^{\text{Logit}}}{\partial \theta_k} = (f_{\theta_k}(\mathbf{X} - \mathbf{y})\mathbf{x}^{(k)})$$

Hence the gradient of the function's curve for any vector of logistic parameters  $\boldsymbol{\theta}$  can be described as:

$$\nabla = \begin{bmatrix} \frac{\partial Q^{\text{Logit}}(\boldsymbol{\theta})}{\partial \theta_1} \\ \frac{\partial Q^{\text{Logit}}(\boldsymbol{\theta})}{\partial \theta_2} \\ \vdots \\ \frac{\partial Q^{\text{Logit}}(\boldsymbol{\theta})}{\partial \theta_k} \end{bmatrix}$$

# Progression of the descent algorithm





# Learning rate

How do we know how much to adjust the parameters by?


- ▶ Intuitively, when our gradient is large we want to make big adjustments because we are far away from the minimum
- ▶ As we get closer, i.e. the gradient is smaller, we want to finetune our adjustments and make smaller steps

So we can scale our adjustment by the size of the gradient

- ▶ Let's call this hyperparameter the **learning rate** ( $\lambda$ )
- ▶  $\theta_{\text{New}} = \theta - \nabla \lambda$

The choice of  $\lambda$  is down to the researcher:

- ▶ Overly-large values will prevent minimisation
- ▶ Overly small values may take too long, or risk converging on *local* minima



# Stochastic gradient descent

Gradient descent can be **expensive**:

- ▶ We have to evaluate all rows in our training data before making any updates to the parameters
- ▶ If we have lots of observations
  1. Each calculation takes a long time
  2. Take many iterations to optimise
- ▶ Instead we can use **stochastic gradient descent** (SGD)
  - ▶ Inspect the loss of each observation (or a random subset) individually
  - ▶ Update the coefficients based on each observation

# Stochastic gradient descent

Under GD, for each iteration:

$$\theta_k \leftarrow \theta_k + \lambda \sum_{i=1}^N (y_i - f_{\theta_k}(\mathbf{x}_i)) \mathbf{x}_i^{(k)}$$

Under SGD, for each iteration:

$$\theta_k \leftarrow \sum_{i=1}^N \theta_k + \lambda (y_i - f_{\theta_k}(\mathbf{x}_i)) \mathbf{x}_i^{(k)}$$

- ▶ SGD typically converges a lot faster than GD
  - ▶ Every iteration we make  $N$  small changes to the parameter estimate
  - ▶ Computationally more efficient (we'll cover this more later in the week)
  - ▶ At the cost of some additional noise in the optimisation process

Coding workshop: writing our own logistic  
regression classifier